DigitalHouse > Coding School

DESAFÍO 3:
Predicción de clicks

Gastón Domínguez

Objetivo

- Elaborar un predictor de clicks
- Aplicación de modelos de clasificación
- Considerar la probabilidad de la predicción

GENERALIDADES

Dataset - Archivos

ctr_15.zip	1,239,524
ctr_16.zip	1,092,413
ctr_17.zip	1,093,607
ctr_18.zip	1,013,581
ctr_19.zip	1,326,180
ctr_20.zip	1,356,066
ctr_21.zip	1,523,221
TOTAL	**8,644,592**

TRAIN

ctr_test_labeled.zip	1,139,639

TEST

Dataset – Features

(#: Features Hasheados)

51 Features

Label	• Etiqueta click (1) o no click (0)
Nums. (6)	• ['auction_age','auction_bidfloor','auction_time', • 'creative_height','creative_width','timezone_offset']
String (41) - #	• ['action_categorical_xx','auction_boolean_xx','auction_categorical_xx', • 'creative_categorical_xx','device_id','device_id_type','gender']
Listas (3) - #	• ['action_list_1','action_list_2','auction_list_0']
Bool (1)	• ['has_video']

APLICACIÓN

Pre-Procesamiento

- Cálculo de media y desvío
  - Para el total del dataset
  - Normalización de variables numéricas

- Reducción del dataset
  - a) Raw Data
  - b) Random Sample (p=25%) -> adoptada
  - c) DownSample (Labels=0)

Pipeline

Extractor

- Eliminación de features de listas
- Estandarización
- Reemplazo de NaNs

Vectorizer

- Feature Hasher
- Combinaciones entre features (n_degree)

Classifier

- SGDClassifier (partial_fit, con batches)
- Regresión logística
- Score: r^2 / log-loss

Grid

Features	• [[:15] , [:]]
N_features	• [2^20, 2^21, 2^22]
degree	• [1 o 2]
l1_ratio	• [0.15, 0.5, 0.85]
alpha	• [0.0001, 0.001, 0.01, 0.1]

TRAIN

TRAIN – Modelos fitteados

features	n_features	degree	l1_ratio	alpha	score
[:15]	2^20	1	0.15	0.0001	0.999
				0.001	0.945
				0.01	0.312
				0.1	-0.044
			0.5	0.0001	0.998
				0.001	0.945
				0.01	0.311
				0.1	-0.044
			0.85	0.0001	0.999
				0.001	0.944
				0.01	0.311
				0.1	-0.044
		2	0.15	0.0001	0.961
				0.001	0.997
				0.01	0.923
				0.1	0.402
			0.5	0.0001	0.980
				0.001	0.988
				0.01	0.923
				0.1	0.352
			0.85	0.0001	0.989
				0.001	0.997
				0.01	0.923
				0.1	0.384

TRAIN – Modelos fitteados

features	n_features	degree	l1_ratio	alpha	score
[:15]	2^20	1	0.15	0.0001	0.999
				0.001	0.945
				0.01	0.312
				0.1	-0.044
			0.5	0.0001	0.998
				0.001	0.945
				0.01	0.311
				0.1	-0.044
			0.85	0.0001	0.999
				0.001	0.944
				0.01	0.311
				0.1	-0.044
		2	0.15	0.0001	0.961
				0.001	0.997
				0.01	0.923
				0.1	0.402
			0.5	0.0001	0.980
				0.001	0.988
				0.01	0.923
				0.1	0.352
			0.85	0.0001	0.989
				0.001	0.997
				0.01	0.923
				0.1	0.384

score

score
0.999
0.945
0.312
-0.044
0.998
0.945
0.311
-0.044
0.999
0.944
0.311
-0.044

TRAIN – Modelos fitteados

features	n_features	degree	l1_ratio	alpha	score
[:15]	2^20	1	0.15	0.0001	0.999
				0.001	0.945
				0.01	0.312
				0.1	-0.044
			0.5	0.0001	0.998
				0.001	0.945
				0.01	0.311
				0.1	-0.044
			0.85	0.0001	0.999
				0.001	0.944
				0.01	0.311
				0.1	-0.044
		2	0.15	0.0001	0.961
				0.001	0.997
				0.01	0.923
				0.1	0.402
			0.5	0.0001	0.980
				0.001	0.988
				0.01	0.923
				0.1	0.352
			0.85	0.0001	0.989
				0.001	0.997
				0.01	0.923
				0.1	0.384

score

score
0.999
0.945
0.312
-0.044
0.998
0.945
0.311
-0.044
0.999
0.944
0.311
-0.044

[:15]	2^20	1	0.5	0.0001	0.998

TEST

TEST – (Label = 0)

Probabilidad de click

Mean (y_proba | Label=0) = 0.13436

Min (y_proba | Label=0) = 0.00068

Max (y_proba | Label=0) = 0.99358

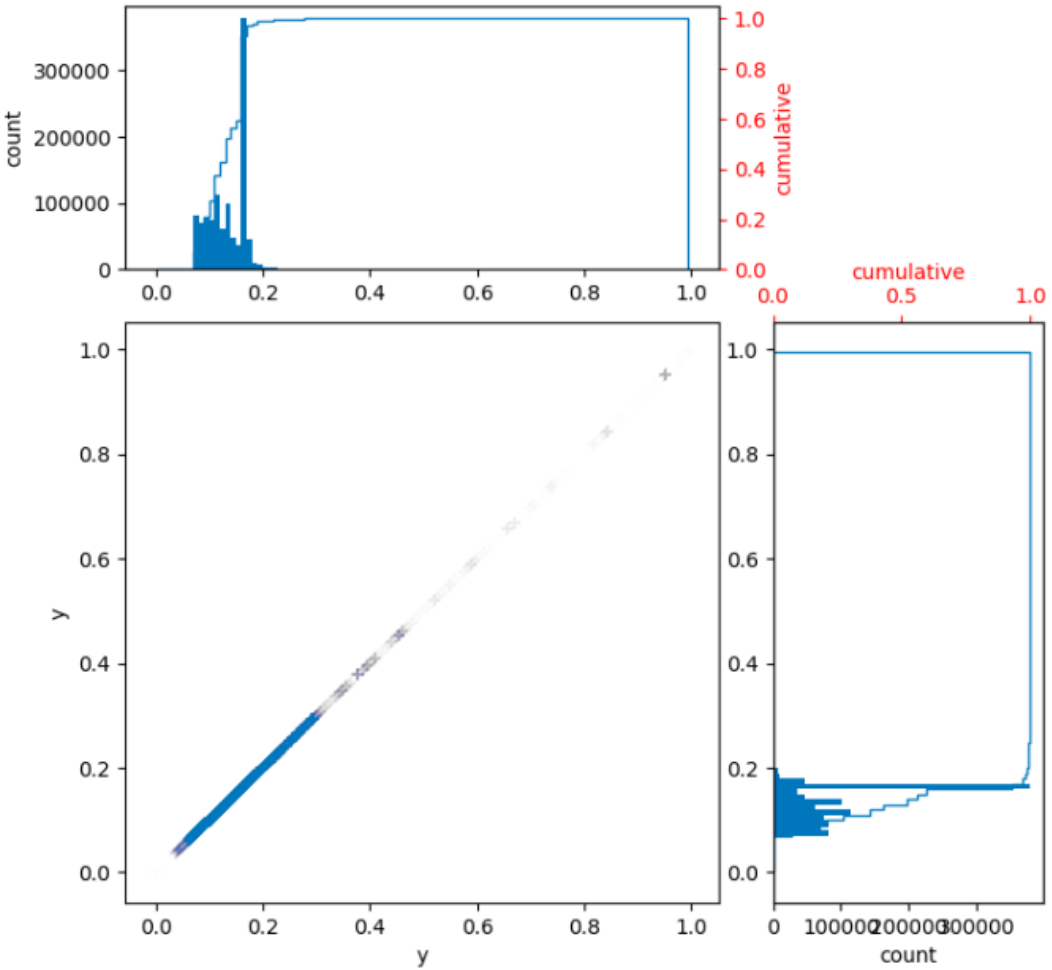

TEST – (Label = 1)

Probabilidad de click

Mean (y_proba | Label=0) = 0.15948

Min (y_proba | Label=0) = 0.05866

Max (y_proba | Label=0) = 0.99885

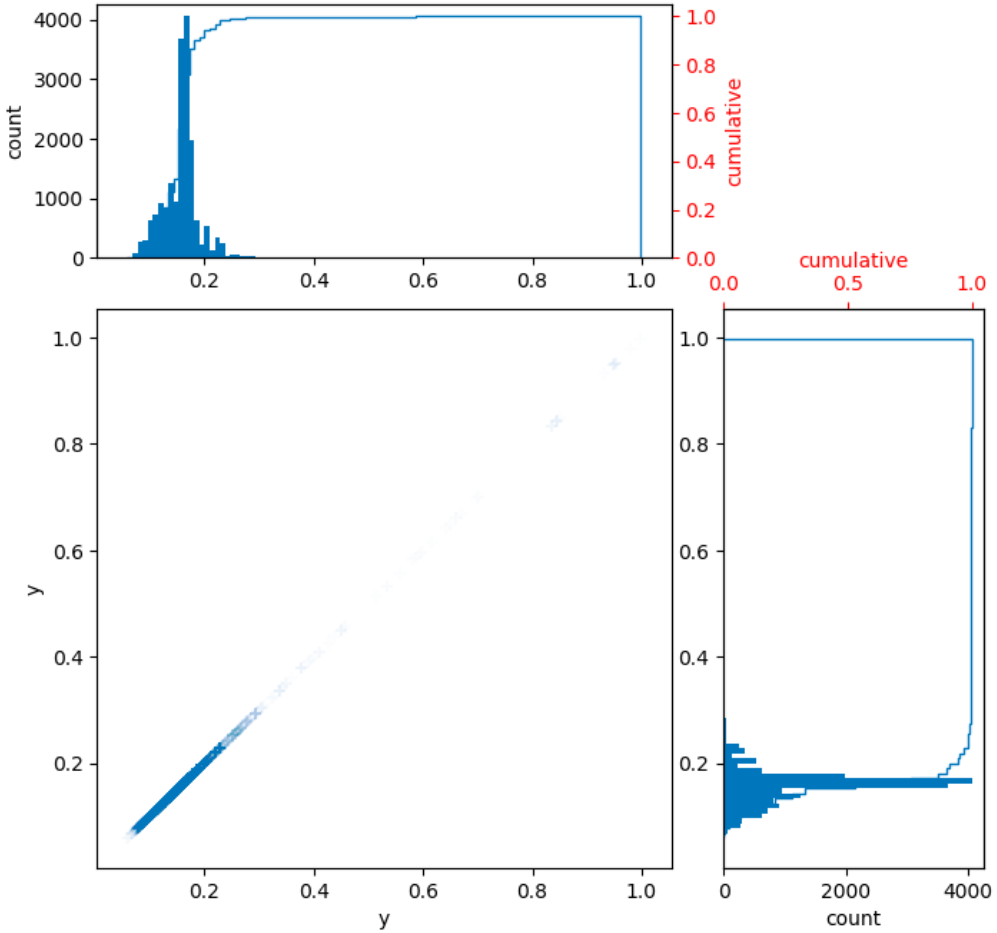

TEST – Label = [0 , 1]

Label = 0

Mean = 0.13436

Min = 0.00068

Max = 0.99358

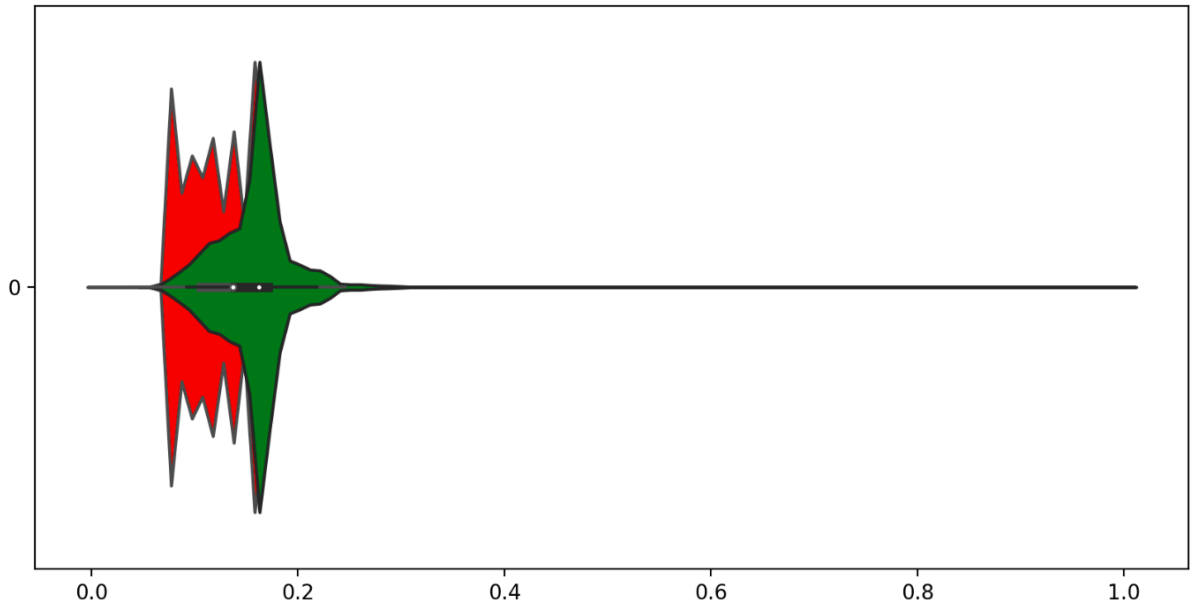

Label = 1

Mean = 0.15948

Min = 0.05866

Max = 0.99885

TEST – Report

	precision	recall	f1-score	support
0	0.98413	0.99971	0.99186	1121523
1	0.11233	0.00226	0.00444	18116
accuracy			0.98386	1139639
macro avg	0.54823	0.50099	0.49815	1139639
weighted avg	0.97028	0.98386	0.97617	1139639

Muchas Gracias!

Preguntas ?